




Scalable blockchains as data layers



Prehistory: meta-protocols

- Mastercoin: “meta-protocol” on top of the Bitcoin blockchain
- Uses the Bitcoin blockchain as a data store, NOT as a state execution engine
- Anyone can compute the Mastercoin state as follows:
 - Let $t[1] \dots t[n]$ be an ordered list of all transactions in the Bitcoin blockchain
 - Let G be the genesis state, and STF be the Mastercoin state transition function
 - $S[1] = STF(G, t[1])$
 - $S[2] = STF(S[1], t[2]) \dots$
 - $S = S[n] = STF(S[n-1], t[n])$




Problems

- Not light client friendly
- MSC activity cannot influence any state outside of the MSC meta-protocol (eg. cannot handle bitcoins)

ZK Rollup

- Contract stores state as Merkle root of `{account ID => (pubkey, balance)}`
- Users send txs of ~13 bytes each plus signature

From	To	Value	Fee	Nonce	Signature
3 bytes	3 bytes	4 bytes	1 byte	2 bytes	32-96 bytes



- Relayer gathers set of transactions `t[1]...t[n]`, creates ZK SNARK:
 - `STF(PRE_STATE, t[1] ... t[n]) = POST_STATE`
 - Each `t[i]` has valid signature
 - `root(PRE_STATE) = r1`
 - `root(POST_STATE) = r2`
- Relayer publishes `r1, r2, t[1]...t[n]` without signatures and SNARK. Contract verifies SNARK and that `root == r1`, and sets `root = r2`

ZK Rollup

- Cost of a rolled-up transaction: ~500k gas for a STARK, plus ~884 gas ($68 * 18$) per transaction
- Note: this scheme can process instant deposits/withdrawals and allows anyone to become a relayer, because unlike Plasma it is NOT susceptible to data unavailability attacks
 - Publishing data required to allow anyone to reconstruct the current balance tree
- Takeaway from Stanford eth1.x workshop: a gas cost reduction in the cost of data (now 68/byte) is likely!

Taking rollup further

- Support more complex state transition functions
- Multiple tokens
- Privacy-preserving computation
- All using the “SNARK + publish tree deltas” paradigm

ZK ZK Rollup

- Users publish txs with SNARKs saying “I have a valid spend certificate for some coin hash in the state. Here is a new coin hash”
 - Eg. user secret: s , coin hash: $h(s + 1)$, spend certificate: $h(s + 2)$
- Transaction does not reveal which coin it is spending. However, each spend certificate can only be used once...
- Relayer only publishes a receipt


Spend certs	New coin values	Fee	New hash
h1: 32 bytes	o1.value: 4 bytes	1 byte	32 bytes
h2: 32 bytes	o2.value: 4 bytes		

Only one hash required for N outputs because it's the Merkle root of all the output coin hashes; the spend would need to provide the Merkle branch to the relayer

- Receipt sufficient to allow anyone to construct updated state
- Gas cost: $105 * 68 = 7140$ per tx


Beacon chain phase 1

- Shard chains as data-only chains
- ~2.8 MB/sec of data availability
- Let's throw things on there so our layer 2's don't have data availability issues!



Eth2 in eth1 light client

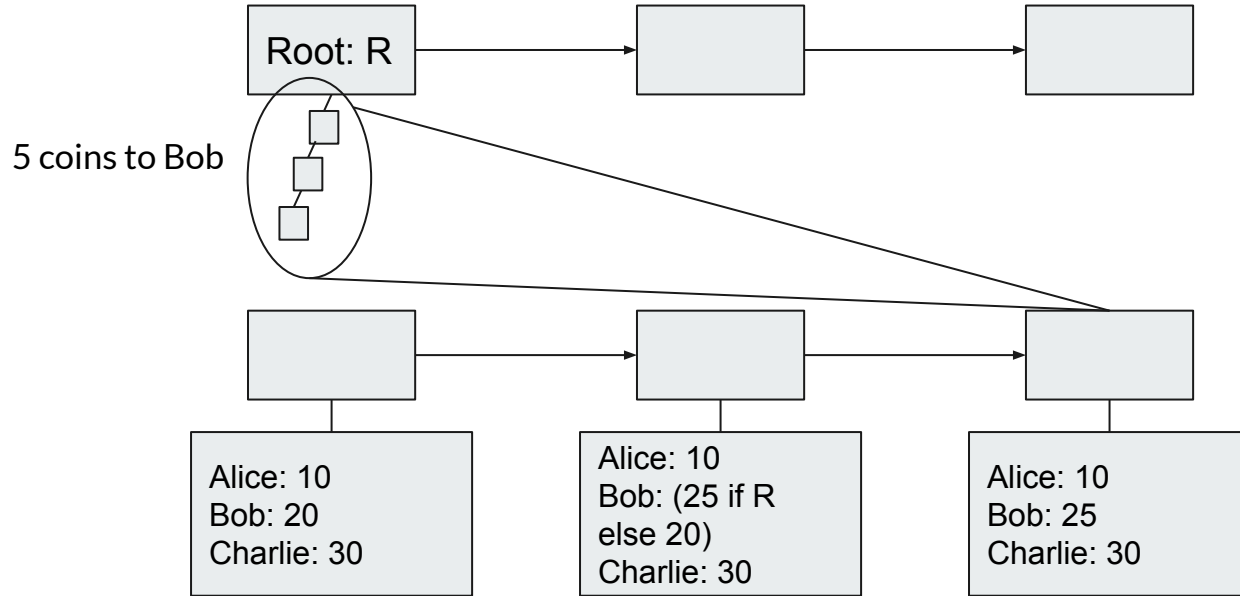
- Light client only requires ~80kB of Merkle branches per 9 days, plus ~500 bytes per header
- Main challenge: need BLS-12-381 precompile in eth1



Other things you can do with scalable data availability engines

- Plasma chains with much more frequent commitments
- Dapps storing messages on-chain
- Blockchain protocols with independent (“sovereign”) state transition functions piggybacking on ethereum for data availability

Speeding up cross-shard transactions



- Store conditional objects in the state of each shard
- Resolve them as we later learn about the actual state roots and branches
- See: <https://ethresear.ch/t/a-layer-2-computing-model-using-optimistic-state-roots/4481>

General-purpose privacy

ZEXE (Zero knowledge EXEcution)

Ledger-Based System

- Supports offline computations
- Provides publicly-verifiable transactions that attest to the correctness of these offline executions

- **Privacy**

- A transaction reveals no information about the offline computation
- Except an upper bound on the number of consumed inputs and created outputs

*Prior work achieves
data privacy but not
function privacy*



Graphic: <https://www.esd.com/news/30/Water-components-leaking-vast-amounts-through-pipeline-its-shotright-focus.com/>

Benefits of the “layer 2 computation” paradigm

- Layer 1 does NOT need to over-complexify to optimize properties
- Layer 2's can upgrade over time, less need for active governance of layer 1
- No need to do “ethereum 3.0 super-quadratic sharding”, can just keep using the beacon chain and increase shard count over time as tech improves
 - “Beacon chain as end of history” thesis